

Regular Expressions

```
library(stringi)
library(stringr)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

Regular expression also known as **Regex** is a tool for describing patterns in strings. The package associated with regex in R are stringr and stringi. Regex is also commonly used in other languages with perhaps slight syntactic changes.

stringr: character manipulation. *stringi*: character string processing facilities.

1) stringi package

`str_detect()` helps detect string patterns in a string. Ensure you have installed stringi and stringr packages.

```
city <- "Los Angeles"
str_detect(city, "os")
```

```
## [1] TRUE
```

```
cities <- c("Los Angeles", "New York", "Atlanta", "New Delhi")
str_detect(cities, "New")
```

```
## [1] FALSE TRUE FALSE TRUE
```

```
# case sensitive
str_detect(cities, "new")
```

```
## [1] FALSE FALSE FALSE FALSE
```

`str_extract()` helps extract a substring from a string and `str_locate()` helps with the start and end points of a pattern.

```
str_extract(cities, "New")
```

```
## [1] NA      "New" NA      "New"
```

```
str_locate(cities, "New")
```

```
##      start end
## [1,]    NA  NA
## [2,]     1   3
## [3,]    NA  NA
## [4,]     1   3
```

`str_sub()` returns a substring of a string given start and end position.

```
# Atlanta has only 7 characters
# thus it returns "nta"
str_sub(cities, start = 5, end = 8)
```

```
## [1] "Ange" "York" "nta"  "Delh"
```

2) Regular Expressions

2.1) Anchors

Anchors help us assert the position, i.e., the beginning or end of the string.

Anchor	Description	Example
<code>^</code>	Matches any string starting with a substring	<code>^New</code>
<code>\$</code>	Matches any string ending with a substring	<code>y\$</code>
Exact Match <code>^\$</code>	Matches the string that starts and ends with substring	<code>^Hi There\$</code>

2.2) Character Classes

Character classes match any characters given a class

Character Class	Match
<code>[aeiou]</code>	Matches vowels
<code>[0-9]</code>	Matches digits
<code>[a-z]</code>	Matches lower case letters
<code>[A-Z]</code>	Matches upper case letters
<code>[a-zA-Z]</code>	Matches both lower and upper case letters
<code>[a-g]</code>	Matches characters a through g

2.3) Quantifiers

Quantifiers quantify the number of instances of a character, group or character class. Your quantifier should be placed after the character/group/character class that is being quantified.

Quantifier	Description
P*	0 or more instances of P
P+	1 or more instances of P
P?	0 or 1 instance of P
P{m}	Exactly m instances of P
P{m,}	At least m instances of P
P{m,n}	Between m and n instances of P

2.4) Logical Operators

Operator	Usage	Example	Explanation
^	Not	[^A-Za-z0-9]	Identifies characters that are not alphanumeric
\	Or	(Apple\ Orange)	Identifies cases that have Apple or Oranges or both

```
str_extract(cities, regex("ew"))
```

```
## [1] NA      "ew" NA      "ew"
```

```
str_extract(cities, regex("^New [a-zA-Z]*"))
```

```
## [1] NA          "New York" NA          "New Delhi"
```

```
banks = c("Bank of America", "Bank of the West", "Citibank",  
          "TD Bank", "Bank of England", "People\'s United Bank")
```

```
#banks starting with "Bank"
```

```
str_detect(banks, regex("^Bank"))
```

```
## [1] TRUE TRUE FALSE FALSE TRUE FALSE
```

```
#banks ending with "Bank" (note: citibank is False as b is lowercase)
```

```
str_detect(banks, regex("Bank$"))
```

```
## [1] FALSE FALSE FALSE TRUE FALSE TRUE
```

```
#bank called "TD Bank" (exact match)
```

```
str_detect(banks, regex("^TD Bank$"))
```

```
## [1] FALSE FALSE FALSE TRUE FALSE FALSE
```

```
# same as above
str_detect(banks, "TD Bank")
```

```
## [1] FALSE FALSE FALSE TRUE FALSE FALSE
```

```
#Note the whitespace after Z in the regular expression
str_extract(banks, regex("^Bank[a-zA-Z ]*"))
```

```
## [1] "Bank of America" "Bank of the West" NA NA
## [5] "Bank of England" NA
```

```
#Without the whitespace after Z, it will not give us the desired output
str_extract(banks, regex("^Bank[a-zA-Z]*"))
```

```
## [1] "Bank" "Bank" NA NA "Bank" NA
```

```
str_extract(banks, regex("^[a-zA-Z ]{1,9}$"))
```

```
## [1] NA NA "Citibank" "TD Bank" NA NA
```

```
print(banks)
```

```
## [1] "Bank of America" "Bank of the West" "Citibank"
## [4] "TD Bank" "Bank of England" "People's United Bank"
```

```
#Detect all bank names containing Bank or bank
str_detect(banks, regex("Bank|bank"))
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE
```

```
#Detect all bank names containg special characters excluding whitespace
str_detect(banks, regex("[^A-Za-z0-9 ]"))
```

```
## [1] FALSE FALSE FALSE FALSE FALSE TRUE
```

```
#List the names of first 10 nba players
# with college name starting with "University"
nba = read.csv("nba2018-players.csv")
```

```
nba %>%
  filter(str_detect(college, regex("^University"))) %>% head(10)
```

```
##           player team position height weight age experience
## 1  Alfonzo McKinnie TOR      SF      80    215  25          0
## 2    Delon Wright TOR      PG      77    183  25          2
## 3   DeMar DeRozan TOR      SG      79    221  28          8
## 4   Jakob Poeltl TOR       C      84    248  22          1
## 5    Nigel Hayes TOR      SF      80    254  23          0
```

## 6	Norman Powell	TOR	SG	76	215	24	2	
## 7	Al Horford	BOS	C	82	245	31	10	
## 8	Jabari Bird	BOS	SG	78	197	23	0	
## 9	Jaylen Brown	BOS	SG	79	225	21	1	
## 10	Kadeem Allen	BOS	PG	75	192	25	0	
##		college		salary	games	minutes	points3 points2	
## 1	University of Wisconsin-Green Bay			815615	14	53	3 5	
## 2	University of Utah			1645200	69	1433	56 145	
## 3	University of Southern California			27739975	80	2711	89 556	
## 4	University of Utah			2825640	82	1524	1 252	
## 5	University of Wisconsin			92160	2	6	2 0	
## 6	University of California, Los Angeles			1471382	70	1062	53 97	
## 7	University of Florida			27734405	72	2277	97 271	
## 8	University of California			0	13	115	3 12	
## 9	University of California			4956480	70	2152	121 252	
## 10	University of Arizona			0	18	107	0 6	
##	points1	points	rebounds	assists	steals	blocks	turnovers	fouls
## 1	2	21	7	1	1	1	3	8
## 2	97	555	198	200	72	33	78	81
## 3	461	1840	315	417	85	22	175	151
## 4	60	567	393	57	39	100	85	212
## 5	0	6	0	0	0	0	1	0
## 6	32	385	119	89	37	16	66	111
## 7	94	927	530	339	43	78	132	138
## 8	6	39	19	8	3	1	8	7
## 9	150	1017	346	114	70	26	124	181
## 10	7	19	11	12	3	2	9	15

```

# List the names of colleges of nba players containing
# special characters (don't consider whitespace and ,
# as special character here)
nba %>%
  filter(str_detect(college, regex("[^A-Za-z, ]"))) %>%
  select(college)

```

##	college
## 1	University of Wisconsin-Green Bay
## 2	Indiana University-Purdue University Indianapolis
## 3	Texas A&M University
## 4	Saint Mary's College of California
## 5	Saint Joseph's University
## 6	Saint Joseph's University
## 7	Saint Joseph's University
## 8	Alabama - Huntsville
## 9	St. John's University
## 10	University of Illinois at Urbana-Champaign
## 11	University of Illinois at Urbana-Champaign
## 12	Saint Mary's College of California
## 13	Texas A&M University
## 14	Texas A&M University
## 15	St. John's University
## 16	Texas A&M University
## 17	Texas A&M University

3) Using Apply

Apply functions are useful when you want to apply a certain operation to all the rows of a list or dataframe. In the example below, we will only consider the case of a data frame. Note that this idea would be useful for your next lab. Here we apply `str_locate()` to locate area code from phone numbers. Using the output of `str_locate()`, we could then extract the area code.

```
phone_num = c(
  "401-501-1111",
  "(401)501-1111",
  "401 501 1111",
  "401-5011111",
  "+408-501-1111")

code_pos = lapply(
  phone_num,
  function(x) str_locate(x, pattern = regex('[0-9]{3}[- ]*'))))

# Use substring() to extract the area code by sharing the start
# and end position with the function as arguments.
code_pos = do.call(rbind, code_pos)
substring(phone_num, code_pos[,1], code_pos[,2]-1)
```

```
## [1] "401" "401" "401" "401" "408"
```

Practice Problems

1. Print only those cities that start with "New" from `city`. Hint: Using `str_detect()`, we could create a mask (a vector of TRUE and FALSE) that could then be subsetting to get the city names starting with "New".

```
cities[str_detect(cities, "New")]
```

```
## [1] "New York" "New Delhi"
```

2. Using `str_match()` explained in cheatsheet, check if `city` has any matches for Los Angeles.

```
str_match(cities, "Los Angeles")
```

```
##      [,1]
## [1,] "Los Angeles"
## [2,] NA
## [3,] NA
## [4,] NA
```

3. Using `nba` dataframe, print names of Players containing Marcus

```
# your code
nba %>% filter(str_detect(player, regex("Marcus"))) %>% select(player)
```

```
##           player
## 1      Marcus Morris
## 2      Marcus Smart
## 3      Marcus Paige
## 4    DeMarcus Cousins
## 5    LaMarcus Aldridge
## 6  Marcus Georges-Hunt
```

4. Using nba dataframe, print names of Universities that contain California or Los Angeles in them.

```
# your code
nba %>% filter(str_detect(college, regex('California|Los Angeles')) %>% select(college)
```

```
##           college
## 1      University of Southern California
## 2      University of California, Los Angeles
## 3      University of California
## 4      University of California
## 5  University of North Carolina, University of California, Los Angeles
## 6      University of California, Los Angeles
## 7      University of California, Los Angeles
## 8      University of California, Los Angeles
## 9      University of California, Los Angeles
## 10     California State University, Fresno
## 11     Saint Mary's College of California
## 12     University of California, Los Angeles
## 13     California State University, Long Beach
## 14     University of California
## 15     California Polytechnic State University, San Luis Obispo
## 16     University of California, Los Angeles
## 17     University of California, Los Angeles
## 18     University of Southern California
## 19     University of Southern California
## 20     California State University, Fullerton
## 21     University of California, Los Angeles
## 22     University of California
## 23     University of California, Los Angeles
## 24     University of California, Los Angeles
## 25     University of Southern California
## 26     California State University, Fresno
## 27     University of California, Los Angeles
## 28     University of California, Los Angeles
## 29     University of California, Los Angeles
## 30     Saint Mary's College of California
## 31     University of Southern California
## 32     University of California
## 33     University of California, Los Angeles
## 34  University of North Carolina, University of California, Los Angeles
## 35     University of California
## 36     University of California, Santa Barbara
```

5. Consider the variable below myVar. Using str_detect and regular expression, detect those strings that contain at least one z and a maximum of three z. Your output should be TRUE TRUE TRUE FALSE

```
myVar = c("bizarre", "bizzarre", "bizzzarre", "bizzzzare")
```

```
# your code
str_detect(myVar, regex("[a-y]+z{1,3}[a-y]+"))
```

```
## [1] TRUE TRUE TRUE FALSE
```

6. Consider the output of ls statement from command line. We have stored it in `file_names`, a vector containing file names. Detect file names that start with STAT154 and have extension .csv
Expected Output: TRUE FALSE FALSE TRUE FALSE TRUE

```
file_names = c(
  "STAT154nba.csv",
  "test0102_STAT154.csv",
  "STAT154myfile.csv.tmp",
  "STAT154_lab.csv",
  "STAT154_HW.pdf",
  "STAT154_test0102.csv")
```

```
# your code
str_detect(file_names, regex("(STAT154)[a-zA-Z0-9_]+(.csv)$"))
```

```
## [1] TRUE FALSE FALSE TRUE FALSE TRUE
```

7. For the variable `newVar`, identify strings that start with a number. Your output should be TRUE FALSE TRUE TRUE FALSE.

```
newVar = c(
  "1 Student(s)",
  "None but 1 Student(s)",
  "5 Students",
  "120! Students",
  "Two Students")
```

```
# your code
str_detect(newVar, regex("[0-9]"))
```

```
## [1] TRUE FALSE TRUE TRUE FALSE
```

8. Detect phone numbers from a given list. Note that these won't have international codes. Your output for given test case should be TRUE TRUE TRUE TRUE FALSE FALSE Hint: The last string in the vector has more than 10 digits.

```
phone_num = c(
  "401-501-1111",
  "(401)501-1111",
  "401 501 1111",
  "4015011111",
  "+408-501-1111",
  "40850211111")
```



```
# your code
str_detect(phone_num, regex("[0-5()-[ ]{6})([1]{4})"))
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE
```

9. In the solution you gave above, can you think of cases where your expression would fail to detect an incorrect phone number? If yes, how could you improve it?

```
# No code, but some text
#
# It would fail to detect an incorrect phone number
# if the character length is longer or shorter than your average 10 character numbers.
```

10. In column of a dataframe, we have stored strings that tell us if someone likes orange juice, apple juice etc. Extract the names of students who like orange juice or oranges. Expected Output: Annie Harry
Hint: Use `str_detect` before using `str_extract()`

```
myVar = c(
  "Annie likes Orange juice", "Sonny likes Apple juice",
  "Katy dislikes Orange juice", "Harry likes Oranges",
  "Charlie likes Apple juice", "Margo like Orange Pie")

df = data.frame(myVar)
```

```
# your code
filtered <- df %>% filter(str_detect(myVar, regex(" likes Orange juice|( likes Oranges)$")))
str_extract(filtered$myVar, regex("^[A-Za-z]+"))
```

```
## [1] "Annie" "Harry"
```

11. Using the idea above, from the `nba` dataframe, pull the last names of all NBA players. Find the frequencies of all the last names. *Hint:* Use `str_locate()` and `substring()`. Eg: `str_locate(nba$player, ...)`. Alternatively, you may choose to try this with `lapply()` but this is slightly tricky.

```
# your code
head(str_sub(nba$player, str_locate(nba$player, pattern = regex('[^ ]+$'))), 10)
```

```
## [1] "McKinnie" "Miles" "Wright" "DeRozan" "VanVleet"
## [6] "Poeltl" "Valanciunas" "Lowry" "Brown" "Nogueira"
```

As the name suggests, it helps look around the string. Look **Arounds** indicate positions just like anchors, `$`, `^`, that we learnt in previous lab.

4) Look Aheads

The expression `A(?=B)` for look-ahead means “look for A, but match only if followed by B”. There may be any pattern instead of A and B.

4.1) Summary of Look-ahead

Type	Syntax	Description
Positive Look Ahead	(?=pattern)	Lookahead Asserts that what immediately follows the current position in the string is pattern
Negative Look Ahead	(?!pattern)	Negative Lookahead Asserts that what immediately follows the current position in the string is not pattern

```
myVar2 = "7 days of 10mm rainfall in Ohio and 2 nights of 15mm rainfall in NYC"
```

```
# extract all digits followed by "mm"
str_extract_all(myVar2, regex('[0-9]+(?=mm)'))
```

```
## [[1]]
## [1] "10" "15"
```

```
# extract all digits not followed by "mm"
str_extract_all(myVar2, regex('[0-9]+ (?!mm)'))
```

```
## [[1]]
## [1] "7 " "2 "
```

5) Look Behinds

Look Behind allows to match a pattern only if there's something before it. This is contrary to lookahead which allows to assert for “what follows”. The expression (?<=B)A matches A, but only if there's B before it.

5.1) Summary of Look-behind

Type	Syntax	Description
Positive Look Behind	(?<=pattern)	Lookbehind Asserts that what immediately precedes the current position in the string is pattern
Negative Look Behind	(?<!pattern)	Negative Lookbehind Asserts that what immediately precedes the current position in the string is not pattern

```
myVar3 = '1 apple costs $1.50 in USA, 2 apples cost £3 elsewhere and 5€ in France'
```

```
# extract only prices
str_extract_all(myVar3, regex('((?<=[£$]{1})[0-9.]+)|([0-9.]+(?=[€]{1}))'))
```

```
## [[1]]
## [1] "1.50" "3" "5"
```

Prictice Problems

1. In these examples, lets try to extract the second university listed for every player in `sample_college` variable, a subset of the `nba` dataset we used earlier.

```
nba = read.csv("nba2018-players.csv")

sample_college = nba[c(1,4,14,18,37,51,56,245,254:256,274:276,291),'college']
sample_college
```

```
## [1] University of Wisconsin-Green Bay
## [2] University of Southern California
## [3] University of California, Los Angeles
## [4] Northern Illinois University, Iowa State University
## [5] University of Colorado, Northern Illinois University
## [6] University of North Carolina, University of California, Los Angeles
## [7] Duquesne University, University of Arizona
## [8] University of Pittsburgh, University of Nevada, Las Vegas
## [9] University of Nebraska, Syracuse University
## [10] Southeast Missouri State University
## [11] Drexel University, University of Louisville
## [12] California State University, Fullerton
## [13] University of North Carolina
## [14] Virginia Commonwealth University
## [15] University of Memphis, University of Kansas
## 150 Levels:  Alabama - Huntsville ... Xavier University
```

In case our input string is University of Memphis, University of Kansas, we should get output as University of Kansas, i.e., print the second university name listed.

```
str_extract(sample_college, regex("(?<=, ).*University.*"))
```

```
## [1] NA
## [2] NA
## [3] NA
## [4] "Iowa State University"
## [5] "Northern Illinois University"
## [6] "University of California, Los Angeles"
## [7] "University of Arizona"
## [8] "University of Nevada, Las Vegas"
## [9] "Syracuse University"
## [10] NA
## [11] "University of Louisville"
## [12] NA
## [13] NA
## [14] NA
## [15] "University of Kansas"
```

2. Using the expressions developed so far, write a function that returns a dataframe with two columns, last name of player and the name of second university that we extracted in last questions. The input to the function would be a subset of `nba` players. You only need to populate the function definition.

Hint: On running the above chunk, your output should be (for first five lines):

	Last_name	Second_Univ
1	McKinnie	
2	DeRozan	
3	Powell	
4	Nader	Iowa State University
5	Silas	Northern Illinois University

```
myfunc = function(df) {
  Last_name <- str_extract(df$player, regex("(?<= ).*"))
  Second_Univ <- str_extract(df$college, regex("(?<=, ).*University.*"))
  new_df <- data.frame(Last_name, Second_Univ)
  ### Write your code here and modify return statement
  return(head(new_df, 10))
}

myfunc(nba[c(1,4,14,18,37,51,56,245,254:256,274:276,291),])
```

```
##      Last_name                Second_Univ
## 1    McKinnie                <NA>
## 2    DeRozan                <NA>
## 3    Powell                <NA>
## 4      Nader            Iowa State University
## 5      Silas      Northern Illinois University
## 6      Drew University of California, Los Angeles
## 7 McConnell      University of Arizona
## 8      Birch      University of Nevada, Las Vegas
## 9      White      Syracuse University
## 10 Cleveland                <NA>
```

3. In webscrapped content, one could retrieve desired content by looking for tags. In this case we would like to retrieve `firstHeading`.

Example: `<h1 id="firstHeading" class="firstHeading" lang="en">University of California, Berkeley</h1>`

Here **University of Berkeley** is the content we need to retrieve, i.e. our `firstHeading`. First headings are always enclosed between

- (a) `<h1 id="firstHeading" class="firstHeading" lang="en">`
- (b) `</h1>`

Write code to extract all occurrences of First Headings in variable `text`. Use positive look ahead and positive look behind to capture the content between (a) and (b) given above.

Your output should be:

```
[[1]]
[1] "Yoshua Bengio"          "Turing Award"
[3] "University of Manchester" "Chicken soup"
```

```
text = '<h1 id="firstHeading" class="firstHeading" lang="en">University of California, Berkeley</h1><h1
```

```
str_extract_all(text, regex('(?"en">)[a-zA-Z ]*(?=</>'))
```

```
## [[1]]
## [1] "Yoshua Bengio"          "Turing Award"
## [3] "University of Manchester" "Chicken soup"
```

6) Some Metacharacters

Expression	Description
<code>\\d</code>	match any digit (same as <code>[0-9]</code>)
<code>\\s</code>	match any whitespace (space, tab)
<code>\\t</code>	match only tab
<code>\\b</code>	match a word boundary
<code>\\A</code>	match the beginning of input
<code>\\Z</code>	match the end of input

Metacharacters have a special meaning during pattern processing. Literals as we learnt in lectures are actual strings that we match. Lets look at some examples.

```
myVar5 = "7 days of 10mm rainfall in Ohio and 2 nights of 15mm rainfall in NYC"
```

```
str_extract_all(myVar5, regex('\\\\d'))
```

```
## [[1]]
## [1] "7" "1" "0" "2" "1" "5"
```

```
str_extract_all(myVar5, regex('\\\\d\\\\d'))
```

```
## [[1]]
## [1] "10" "15"
```

```
str_extract_all(myVar5, regex('\\\\d+'))
```

```
## [[1]]
## [1] "7" "10" "2" "15"
```

```
str_extract_all(myVar5, regex('\\\\w'))
```

```
## [[1]]
## [1] "7" "d" "a" "y" "s" "o" "f" "1" "0" "m" "m" "r" "a" "i" "n" "f" "a" "l" "l"
## [20] "i" "n" "0" "h" "i" "o" "a" "n" "d" "2" "n" "i" "g" "h" "t" "s" "o" "f" "1"
## [39] "5" "m" "m" "r" "a" "i" "n" "f" "a" "l" "l" "i" "n" "N" "Y" "C"
```

```
str_extract_all(myVar5, regex('\\w+'))
```

```
## [[1]]
## [1] "7"      "days"  "of"      "10mm"    "rainfall" "in"
## [7] "Ohio"   "and"    "2"       "nights"  "of"       "15mm"
## [13] "rainfall" "in"     "NYC"
```

Practice Problems

1. Write code to detect three digit area code The output for this case should be TRUE TRUE FALSE FALSE.
Hint: Replace periods with solution `^[...]?\\d{...}[...]?$`

```
test_case = c('401', '(401)', '4015', '+401')
str_detect(test_case, regex('^([]?\\d{3}[)]?.$'))
```

```
## [1] TRUE TRUE FALSE FALSE
```

2. To the previous expression make changes to detect first 6 numbers. Your output here should be TRUE TRUE TRUE TRUE FALSE FALSE

```
test_case2 = c('401 501', '401-501', '(401)501', '401501', '+401501', '4011501')
str_detect(test_case2, regex('^([]?\\d{3}[ ]-]?\\d{3}$'))
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE
```

3. Now consider the actual problem. Modify solution to previous question to find a pattern for detecting phone numbers. Your output for given test case should be TRUE TRUE TRUE TRUE FALSE FALSE

```
phone_num = c(
  "401-501-1111",
  "(401)501-1111",
  "401 501 1111",
  "4015011111",
  "+408-501-1111",
  "4085021111")
str_detect(phone_num, regex('^([]?\\d{3}[ ]-]?\\d{3}[- ]?\\d{4}$'))
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE
```